

CS 480: Database Systems

Lecture 13

February 13, 2013

Aggregation with GROUP BY

- INSTRUCTOR(name,dept,salary)
- Retrieve the faculty budget for each department (what they pay in total to their professors).

```
SELECT      dept, SUM(salary) AS budget
FROM        instructor
GROUP BY    dept
```

The attributes in the GROUP BY clause have to also appear in the SELECT clause to appear in the result.

Database Modification Operations

- So far all we've done is: accessing the data (querying).
- Now we can also manipulate the data:
 - Inserting new tuples
 - Deleting tuples
 - Updating tuples

Deletion

- Format:

```
DELETE FROM  $r$ 
WHERE       $P$ 
```

- r is the relation from which to choose the tuples to delete.
- P is a predicate that selects the tuples to remove.
- If the WHERE clause is omitted, all tuples in r are deleted.
- Operates on only one relation.

Deletion

- INSTRUCTOR(name,dept,salary)

Deletion

- INSTRUCTOR(name,dept,salary)
- Example #1: Delete all instructors that are part of the CS department.

Deletion

- INSTRUCTOR(name,dept,salary)
- Example #1: Delete all instructors that are part of the CS department.

```
DELETE FROM instructor
WHERE      dept= 'CS'
```

Deletion

- INSTRUCTOR(name,dept,salary)
- Example #1: Delete all instructors that are part of the CS department.

```
DELETE FROM instructor
WHERE      dept= 'CS'
```

- Example #2: Delete all instructors that have a salary between \$13,000 and \$15000.

Deletion

- INSTRUCTOR(name,dept,salary)
- Example #1: Delete all instructors that are part of the CS department.

```
DELETE FROM instructor
WHERE      dept= 'CS'
```

- Example #2: Delete all instructors that have a salary between \$13,000 and \$15000.

```
DELETE FROM instructor
WHERE      salary >= 13000 AND
           salary <= 15000
```

Deletion

- Deletion is from only one relation.
- However, we can reference other relations in SELECT-FROM-WHERE format in nested subqueries inside in the WHERE clause.

Deletion

- STUDENT(id,name,address,GPA)
ENROL(id,course,grade)

Deletion

- STUDENT(id,name,address,GPA)
ENROL(id,course,grade)
- Example #3: Delete the records of every student that got an F in CS480.

Deletion

- STUDENT(id,name,address,GPA)
ENROL(id,course,grade)
- Example #3: Delete the records of every student that got an F in CS480.

```
DELETE FROM student
WHERE      id IN (SELECT id
                  FROM enrol
                  WHERE grade='F' AND
                        course='CS480' )
```

Deletion

- INSTRUCTOR(name,dept,salary)
- Consider the following query:

```
DELETE FROM instructor
WHERE      salary < (SELECT AVG(salary)
                     FROM    instructor)
```

Deletion

- INSTRUCTOR(name,dept,salary)
- Consider the following query:

```
DELETE FROM instructor
WHERE      salary < (SELECT AVG(salary)
                     FROM    instructor)
```

Deletion processing note:

Selection of all tuples to be deleted should be done prior to the actual deletion of the tuples. Why?

Insertion

- Two possible ways:
 1. Specify a tuple to be inserted
 2. Write a query whose result is a set of tuples to be inserted.
- Attribute values for inserted tuples must be members of the corresponding attribute's domain.

Insertion

- Insertion by specifying a tuple to be inserted
- `COURSE(name,title,dept,credits)`

Insertion

- Insertion by specifying a tuple to be inserted
- COURSE(name,title,dept,credits)

```
INSERT INTO course  
VALUES ('CS480', 'Database Systems', 'CS', 4)
```

Insertion

- Insertion by specifying a tuple to be inserted
- `COURSE(name,title,dept,credits)`

```
INSERT INTO course  
VALUES ('CS480', 'Database Systems', 'CS', 4)
```

Order of values should be based in the order that the schema was defined.
If order is unknown then one can specify the order explicitly:

Insertion

- Insertion by specifying a tuple to be inserted
- `COURSE(name,title,dept,credits)`

```
INSERT INTO course  
VALUES ('CS480', 'Database Systems', 'CS', 4)
```

Order of values should be based in the order that the schema was defined.
If order is unknown then one can specify the order explicitly:

```
INSERT INTO course(name,title,dept,credits)  
VALUES ('CS480', 'Database Systems', 'CS', 4)
```

Insertion

- Insertion by specifying a tuple to be inserted
- `COURSE(name,title,dept,credits)`

```
INSERT INTO course  
VALUES ('CS480', 'Database Systems', 'CS', 4)
```

Order of values should be based in the order that the schema was defined.
If order is unknown then one can specify the order explicitly:

```
INSERT INTO course(name,title,dept,credits)  
VALUES ('CS480', 'Database Systems', 'CS', 4)
```

```
INSERT INTO course(name,dept,credits,title)  
VALUES ('CS401', 'CS', 4, 'Algorithms')
```

Insertion

- Insertion by writing a query whose result is a set of tuples to be inserted.

Insertion

- Insertion by writing a query whose result is a set of tuples to be inserted.
- INSTRUCTOR(name,dept,salary)
STUDENT(id,name,address,dept,tot_credit)

Insertion

- Insertion by writing a query whose result is a set of tuples to be inserted.
- INSTRUCTOR(name,dept,salary)
STUDENT(id,name,address,dept,tot_credit)
- Suppose we want to make every CS student an instructor with a salary of \$18,000 if they have earned more than 150 credits.

Insertion

- Insertion by writing a query whose result is a set of tuples to be inserted.
- INSTRUCTOR(name,dept,salary)
STUDENT(id,name,address,dept,tot_credit)
- Suppose we want to make every CS student an instructor with a salary of \$18,000 if they have earned more than 150 credits.

```
INSERT INTO instructor
    SELECT name,dept,18000
    FROM    student
    WHERE   dept='CS' AND tot_credit > 150
```

Updates

- Used for changing a value in a tuple without changing all values in the tuple.
- We can choose which tuples to update.
- Format:

```
UPDATE  r
SET      A1=a1, A2=a2, ..., An=an
WHERE    P
```

- A_i 's are attributes, a_i 's are arithmetic expressions, P is a predicate.

Updates

- INSTRUCTOR(name,dept,salary)
- Annual salary increases need to be made.
All instructor salaries will be increased by 5%.

Updates

- INSTRUCTOR(name,dept,salary)
- Annual salary increases need to be made.
All instructor salaries will be increased by 5%.

```
UPDATE instructor
SET      salary = salary*1.05
```

Updates

- INSTRUCTOR(name,dept,salary)
- Annual salary increases need to be made.
All instructors with salary less than \$70,000 will have their salaries increased by 5%.

Updates

- INSTRUCTOR(name,dept,salary)
- Annual salary increases need to be made. All instructors with salary less than \$70,000 will have their salaries increased by 5%.

```
UPDATE instructor
SET      salary = salary*1.05
WHERE    salary < 70000
```

Updates

- INSTRUCTOR(name,dept,salary)
- All instructors that earn over \$100,000 will get a 3% raise. All other instructors will get a 5% raise.

Updates

- INSTRUCTOR(name,dept,salary)
- All instructors that earn over \$100,000 will get a 3% raise. All other instructors will get a 5% raise.

```
UPDATE instructor
SET      salary = salary*1.03
WHERE    salary > 100000
```

```
UPDATE instructor
SET      salary = salary*1.05
WHERE    salary <= 100000
```


Updates

- INSTRUCTOR(name,dept,salary)
- All instructors that earn over \$100,000 will get a 3% raise. All other instructors will get a 5% raise.

```
UPDATE instructor
SET    salary = CASE
                WHEN salary <= 100000 then salary*1.05
                ELSE salary*1.03
        END
```

Updates with CASE construct

- In general:

```
UPDATE  r
SET     A = CASE
           WHEN P1 then a1
           WHEN P2 then a2
           ...
           WHEN Pn then an
           ELSE a0
        END
```

Updates

- One could also update a value with the result of a scalar query (one that returns just one value).
- STUDENT(id,name,address,tot_credit)
COURSE(name,title,credits)
ENROL(id,course_name,grade)

```
UPDATE student AS s
SET    tot_credit = (
        SELECT SUM(credits)
        FROM    enrol NATURAL JOIN course
        WHERE   s.id=enrol.id AND
                enrol.grade <> 'F' )
```

Data Definition Language

- SQL DDL provides commands for:
 - Defining relation schemas
 - Deleting relations
 - Modifying relation schemas
- When defining a relation in SQL we can specify:
 - The schema for each relation
 - Types of values associated with each attribute
 - Other constraints (integrity, access, storage)

Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.

Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).

Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p,d*)**. Fixed point number, with *p* digits and with *d* of the digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.

Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p,d*)**. Fixed point number, with *p* digits and with *d* of the digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- Many others not covered here

CREATE TABLE construct

- SQL relations are defined using the CREATE TABLE command:

```
CREATE TABLE r
    (A1 D1,
     A2 D2,
     ...
     An Dn,
     <integrity_constraint1>,
     ...
     <integrity_constraintk>)
```

Some Integrity Constraints

- PRIMARY KEY (B_1, B_2, \dots, B_m)
 - B's must be a subset of the already specified attributes (A's).
 - Requires the key attributes to be unique.
 - Optional, but it's generally good practice to specify one.
- FOREIGN KEY (C_1, C_2, \dots, C_d) REFERENCES s
 - The values of these C's should correspond to values of these attributes of some tuple in relation s .
- SQL will prevent any update to the database that violates an integrity constraint.

SQL Beer-drinker Database

Relational Schema:

DRINKER(d-name,d-city), **BAR**(ba-name,ba-city), **BEER**(be-name,type),
FREQUENTS(d-name,ba-name), SERVES(ba-name,be-name),
LIKES(d-name,be-name)

SQL schema definition

```
CREATE TABLE drinker
    (d_name VARCHAR(30),
     d_city VARCHAR(20),
     PRIMARY KEY (d_name));
CREATE TABLE bar
    (ba_name VARCHAR(30),
     ba_city VARCHAR(20),
     PRIMARY KEY (ba_name));
CREATE TABLE beer
    (be_name VARCHAR(30),
     type     VARCHAR(20),
     PRIMARY KEY (be_name));
```

SQL Beer-drinker Database

Relational Schema:

DRINKER(d-name,d-city), BAR(ba-name,ba-city), BEER(be-name,type),
FREQUENTS(d-name,ba-name), **SERVES(ba-name,be-name)**,
LIKES(d-name,be-name)

SQL schema definition

```
CREATE TABLE frequents
    (d_name  VARCHAR(30),
     ba_name VARCHAR(30),
     FOREIGN KEY(d_name) REFERENCES drinker,
     FOREIGN KEY(ba_name) REFERENCES bar);

CREATE TABLE serves
    (ba_name VARCHAR(30),
     be_name VARCHAR(30),
     FOREIGN KEY(be_name) REFERENCES beer,
     FOREIGN KEY(ba_name) REFERENCES bar);

CREATE TABLE likes
    (d_name  VARCHAR(30),
     be_name VARCHAR(30),
     FOREIGN KEY(be_name) REFERENCES beer,
     FOREIGN KEY(d_name) REFERENCES drinker);
```

CHECK clause

- An additional integrity constraint.
- When applied to an attribute declaration it specifies a predicate that must be satisfied by every tuple in the relation.
- Common use is to ensure that attribute values satisfy specified conditions (creates a powerful type system)
- Syntax: `CHECK (P)`

CHECK clause

```
CREATE TABLE beer
    (be_name VARCHAR(30),
     type     VARCHAR(20),
     PRIMARY KEY (be_name),
     CHECK (type IN ('ale', 'lager',
                    'stout', 'pilsner',
                    'porter', 'bitter',
                    'other')));
```

CHECK clause

```
CREATE TABLE employee
(id      CHAR(10)
name    VARCHAR(30),
salary  NUMERIC(10,2),
PRIMARY KEY (id),
CHECK (salary > 0)) ;
```

Deleting Tables

- To remove a relation from an SQL database we use the DROP TABLE command.
- **Syntax:** DROP TABLE r

Altering Tables

- We can use the ALTER TABLE command to change the scheme of a relation by:
 - Adding attributes to the relation
 - Syntax: `ALTER TABLE r ADD A D;`
 - Where *r* is an existing relation, *A* is the name of the new attribute and *D* is the domain of *A*.
 - All values for the new attribute are initially assigned the special value `NULL`.
 - Removing attributes from the relation
 - Syntax: `ALTER TABLE r DROP A;`
 - Where *r* is an existing relation and *A* is an attribute of *r*.

Relational Calculus

Tuple Relational Calculus

- Relational algebra expressions provide a sequence of procedure that generates the answer to our query.
- Tuple Relational Calculus in contrast is a formal query language that is **nonprocedural**.
- It describes the desired information without giving a specific procedure for obtaining the information.

Tuple Relational Calculus

- General format of a relational calculus expression:

$$\{ x \mid \langle \text{condition} \rangle \}$$

Example #1

- student1(STUDENT)
- student2(STUDENT)
- STUDENT = {id,name,major,GPA}
- RA Query: student1 \cup student2
- Calculus:

$$\{ y \mid y \in \text{student1} \vee y \in \text{student2} \}$$

Example #2

- student(STUDENT)
- STUDENT = {id,name,major,GPA}
- RA Query: $\Pi_{id}(\text{student})$
- Calculus:

$$\{ y \mid \exists x \in \text{student} (x[id] = y[id]) \}$$

Example #3

- student(STUDENT), enrol(ENROL)
- STUDENT = {id,name,major,GPA}
- ENROL = {id,course,grade}
- RA Query: $\Pi_{\text{name,grade}}(\text{student} \bowtie \text{enrol})$
- Calculus:

$$\{ x \mid \exists y \in \text{student} \exists z \in \text{enrol} \\ (x[\text{name}] = y[\text{name}] \wedge \\ x[\text{grade}] = z[\text{grade}] \wedge \\ y[\text{id}] = z[\text{id}]) \}$$